

Parallel Computing Lab

Polynomial Multiplication using FFT

By: Dhruv Kohli
Roll No.:120123054
Dept.:M&C

Introduction

Let $A = [a_0 \ a_1 \ a_2 \ \dots \ a_{(n-1)}]$, $B = [b_0 \ b_1 \ b_2 \ \dots \ b_{(n-1)}]$ be the vectors of coefficients of two n -degree-bound polynomials $A(x)$ and $B(x)$ respectively.

Assumption: n is a power of 2

Problem Statement: To compute $C = [c_0 \ c_1 \ c_2 \ \dots \ c_{(2n-2)}]$ which is a vector of coefficients of an $(2n-1)$ -degree-bound polynomial $C(x)$ which is computed as:
$$C(x) = A(x) * B(x)$$

Brute-force polynomial multiplication algorithm has a time complexity of $O(n^2)$.

Another algorithm which makes use of point-wise representation of a polynomial (i.e. every n -degree polynomial can be uniquely represented by $(n+1)$ distinct points $\{(x_i, y_i) | i=1, 2, \dots, n\}$), which we can derive in $O(\lg(n))$ time by interpolating the given polynomial on $2n$ th complex roots** of unity using sequential FFT, and then computes $C(x)$ on those $2n$ th complex roots of unity using $C(x_i) = A(x_i) * B(x_i)$ in $O(n)$ time and finally extrapolates the computed values of C using inverse FFT in $O(\log(n))$ time again to get back the coefficients of C and hence leading to a final time complexity of $O(n \log(n))$ for this algorithm.

Through this assignment, we demonstrate that, given infinite number of processors, we can achieve a time complexity of $O(\lg(n))$ and hence a parallelism of $O(1/n)$ using the so called parallel FFT algorithm for interpolation and extrapolation step. Note that the pointwise multiplication step takes constant time $O(1)$ given infinite number of processors.

**We chose $2n$ th complex roots of unity instead of n because we require $2n-1$ points to represent C pointwise since it is a polynomial of degree $2n-2$.

Parallel FFT Algorithm

```
Recursive_Parallel_FFT(A, IN_OR_EX) {
    n=A.length;
    if(n==1)
        return A;
    w_n=exp(2*IN_OR_EX*pi*i/n);
    w=1;
    A_e <- complex_empty_vector;
    A_o <- complex_empty_vector;
    PARALLEL_THREAD_ARRAY_T
    for(int i = 0; i < n/2; ++i) {
        A_e.push_back(A[2*i]);
        A_o.push_back(A[2*i+1]);
    }
    SYNCH_BARRIER_THREAD_ARRAY_T

    PARALLEL_THREAD_T1 {
        Y_e <- Recursive_Parallel_FFT(A_e)
    }
    PARALLEL_THREAD_T2 {
        Y_o <- Recursive_Parallel_FFT(A_o)
    }
    SYNCH_BARRIER_T1
    SYNCH_BARRIER_T2

    Y <- complex_vector(size=n);
    PARALLEL_THREAD_ARRAY_T
    for(int i = 0; i < n/2; ++i) {
        Y[k] <- Y_e[k] + w*Y_o[k];
        Y[k+n/2] <- Y_r[k] - w*Y_o[k];
        w *= w_n
    }
    SYNCH_BARRIER_THREAD_ARRAY_T
    return Y
}
```

Proposed Algorithm

Fast Polynomial Multiplication:

Let $A_{in} = [a_0 \ a_1 \ a_2 \ \dots \ a_{(n-1)}]$, $B_{in} = [b_0 \ b_1 \ b_2 \ \dots \ b_{(m-1)}]$ and let A and B are complex vectors of same elements as A_{in} and B_{in} but with appropriate number of zeros appended so as to make the size of both A and B equal to $\text{nearUpperPowerOf2}(\text{degree}(C) + 1)$ where $\text{degree}(C) = m+n-2$.

```
Fast_Poly_Multi(A,B) {  
    PARALLEL_THREAD_T1 {  
        A_complex <- Recursive_Parallel_FFT(A, +1);  
    }  
    PARALLEL_THREAD_T2 {  
        B_complex <- Recursive_Parallel_FFT(B, +1);  
    }  
    SYNCH_BARRIER_T1  
    SYNCH_BARRIER_T2  
  
    n = A_complex.length  
    C_complex <- empty complex vector  
  
    PARALLEL_THREAD_ARRAY_T  
    for(int i = 0; i < n; ++i) {  
        C_complex.push_back(A_complex[i]*B_complex[i]);  
    }  
    SYNCH_BARRIER_THREAD_ARRAY_T  
  
    C <- Recursive_Parallel_FFT(A, -1);  
  
    return C.real;    //here C.real[i] = C[i].real  
}
```

Theoretical Analysis

CASE 1: #Proc == 1

Proposed Algo is equivalent to sequential FFT. Hence, $T(1) = O(n \lg(n))$

Case 2: #Proc == infinite

Blue colored region in proposed algorithm takes $O(\lg(n))$ time.

REASON:

Each call to Recursive_Parallel_FFT with an 'n' sized vector results in the formation of a tree of height $\lg(n)$.

Argument: All Vertices at same level can be processed simultaneously (since we have infinite # of processors). Therefore, total time complexity is of order of height of the tree i.e. $O(\lg(n))$.

Orange Colored Region in proposed algorithm takes constant time with infinite # of processors (launch one thread per computation and hence perform the net computation simultaneously)

Hence, $T(\text{infinite}) = O(\lg(n))$

Therefore Parallelism = $T(\text{infinite})/T(1) = O(1/n)$

Case 3: #proc = p

The blue colored region will take time of order $O(n \lg(n)/p)$.

Argument: Assuming that each processor has equal share of computation and since the sequential time complexity was $O(n \lg(n))$, the net time complexity of blue colored region with p processor will be $O(n \lg(n)/p)$.

The orange colored region will take time of order $O(n/p)$.

Hence, $T(p) = O(n \lg(n)/p)$.

openMP implementation

*par_dft is equivalent to Fast_Poly_Multi

```
std::vector<value_type> par_dft(std::vector<value_type> A, std::vector<value_type> B) {
    std::complex<value_type> zero(0,0);
    int degC = A.size() + B.size() - 1;
    int npo2_degC = nearPo2(degC);

    std::vector< std::complex<value_type> > A_c(npo2_degC, zero), B_c(npo2_degC, zero),
    A_ex_c(npo2_degC, zero), B_ex_c(npo2_degC, zero), C_c(npo2_degC, zero),
    C_in_c(npo2_degC, zero);

    #pragma omp parallel for
    for(int i = 0; i < (A.size()); ++i) {
        addElem(&A_c[i], A[i]);
    }
    #pragma omp parallel for
    for(int i = 0; i < (B.size()); ++i) {
        addElem(&B_c[i], B[i]);
    }
    #pragma omp parallel num_threads(2)
    {
        int i = omp_get_thread_num();

        if (i == 0){
            ex_o_in_polate_parallel(&A_c, EXTRAPOLATE, &A_ex_c);
        }
        if (i == 1 || omp_get_num_threads() != 2){
            ex_o_in_polate_parallel(&B_c, EXTRAPOLATE, &B_ex_c);
        }
    }

    //multiply pointwise

    #pragma omp parallel for
    for(int i = 0; i < (A_ex_c.size()); ++i) {
        multiply_kernel(A_ex_c[i], B_ex_c[i], &C_c[i]);
    }

    ex_o_in_polate_parallel(&C_c, INTERPOLATE, &C_in_c);

    std::vector<value_type> res(degC, 0);

    #pragma omp parallel for
    for(int i = 0; i < (degC); ++i) {
        addElem_v(&res[i], C_in_c[i].real()/(1.0*npo2_degC));
    }

    return res;
}
```

***ex_o_in_polate_parallel** is equivalent to Recursive Parallel FFT

```
void ex_o_in_polate_parallel(std::vector< std::complex<value_type> > *A,
                             int TYPE,
                             std::vector< std::complex<value_type> > *A_res) {
    if(A->size() == 1) {
        (*A_res)[0] += (*A)[0];
        return;
    }
    unsigned long N = A->size(); //will always be a power of 2
    std::complex<value_type> w_N(std::cos(TYPE*2.0*M_PI/N), std::sin(TYPE*2.0*M_PI/N));
    std::complex<value_type> w(1.0, 0);
    std::complex<value_type> zero(0,0);
    std::vector< std::complex<value_type> > A_o(N/2, zero), A_e(N/2, zero);

    int cntO = 0, cntE = 0;
    if(omp_get_thread_num() > 1) {
        #pragma omp parallel for
        for(int i = 0; i < N/2; ++i) {
            addElem_c(&A_e[cntE++], (*A)[i*2]);
            addElem_c(&A_o[cntO++], (*A)[i*2+1]);
        }
    } else {
        for(int i = 0; i < N/2; ++i) {
            addElem_c(&A_e[cntE++], (*A)[i*2]);
            addElem_c(&A_o[cntO++], (*A)[i*2+1]);
        }
    }

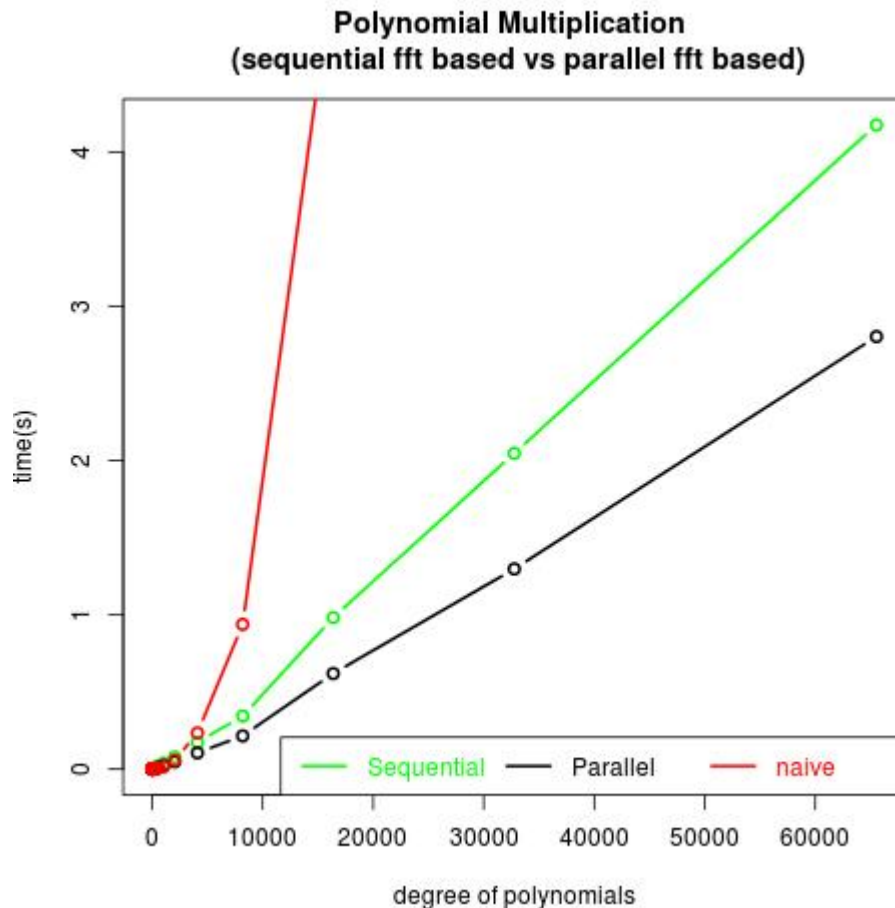
    std::vector< std::complex<value_type> > A_o_c(A_o.size(), zero), A_e_c(A_e.size(), zero);
    #pragma omp parallel num_threads(2)
    {
        int i = omp_get_thread_num();

        if (i == 0){
            ex_o_in_polate_parallel(&A_e, TYPE, &A_e_c);
        }
        if (i == 1 || omp_get_num_threads() != 2){
            ex_o_in_polate_parallel(&A_o, TYPE, &A_o_c);
        }
    }

    if(omp_get_num_threads() > 1) {
        #pragma omp parallel for
        for(int i = 0; i < N/2; ++i) {
            addElem_c(&((*A_res)[i]), A_e_c[i]+w*A_o_c[i]);
            addElem_c(&((*A_res)[i+N/2]), A_e_c[i]-w*A_o_c[i]);
            w = w*w_N;
        }
    } else {
        for(int i = 0; i < N/2; ++i) {
            addElem_c(&((*A_res)[i]), A_e_c[i]+w*A_o_c[i]);
            addElem_c(&((*A_res)[i+N/2]), A_e_c[i]-w*A_o_c[i]);
            w = w*w_N;
        }
    }
}
```

Experimental Analysis and Conclusion

Following graph represents the experimental analysis with # of processors equal to 8.



We conclude that the fast polynomial multiplication algorithm is approximately 1.5 times faster (for degree of polynomial of order $\sim 2^{20}$) and 2 times faster (for degree of polynomial of order $\sim 2^{16}$) than the sequential one in practice on an 8 core machine and we are affirmative that the practical speedup will be much more higher on GPUs with thousand of processors.

NOTE: A C++11 thread based implementation has also been done but probably we are making some conceptual mistakes preventing any speed up even in comparison with sequential dft code.